

CMSC726 MACHINE LEARNING

PROJECT 3

Unsupervised Learning

Angjoo Kanazawa, Ran Liu, and Austin Myers

December 6, 2011

1 PCA and Kernel PCA

1.1 WU1

Depending exactly on your random data, one or more of these lines might not pass exactly through the data as we would like it to. Why not?

As shown in figure 1, the eigenvectors are actually what we would expect, with the first vector accounting for the axis of skew and the second vector orthogonal to the first. This is intuitive since the skew should be the primary source of variance, and the second eigenvector should simply be perpendicular to the first since it's a two dimensional projected space.

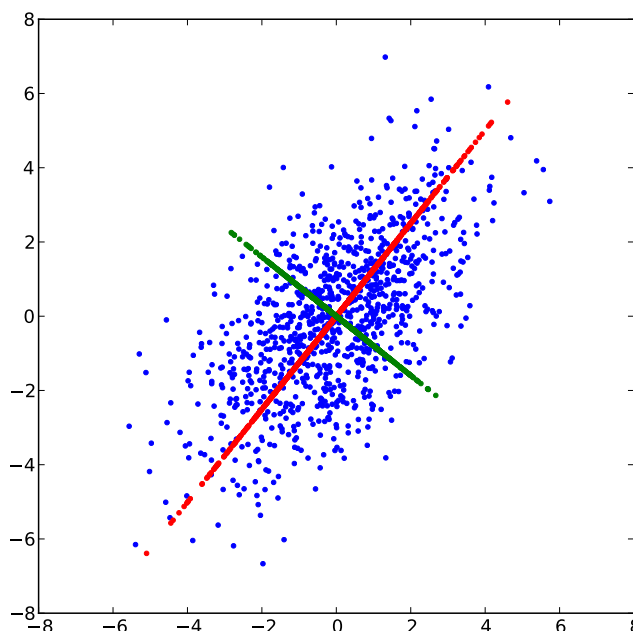


Figure 1: Eigenvectors and projected data

1.2 WU2

Plot the normalized eigenvalues (include the plot in your writeup). How many eigenvectors do you have to include before you've accounted for 90% of the variance? 95%? (Hint: see function `cumsum`.)

We need to include 81 eigenvectors to account for 90% of the variance and 135 eigenvectors to account for 95%. The eigenvalues are shown in

figure 2 and the cumulative sum of the eigenvalues is shown in figure 3 with labels for the two points of interest.

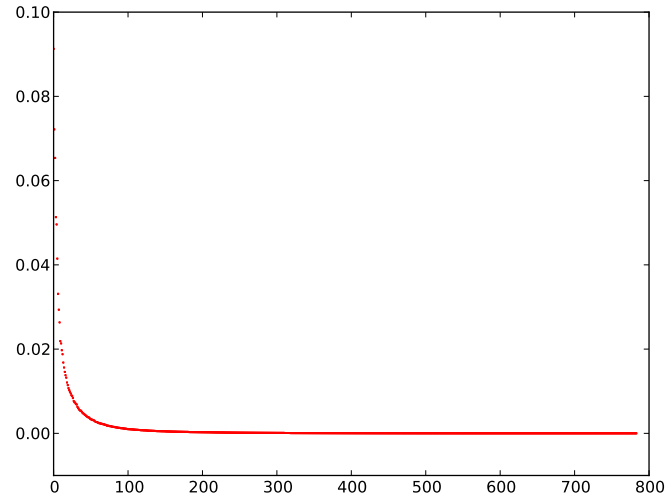


Figure 2: Plot of the normalized eigenvalues

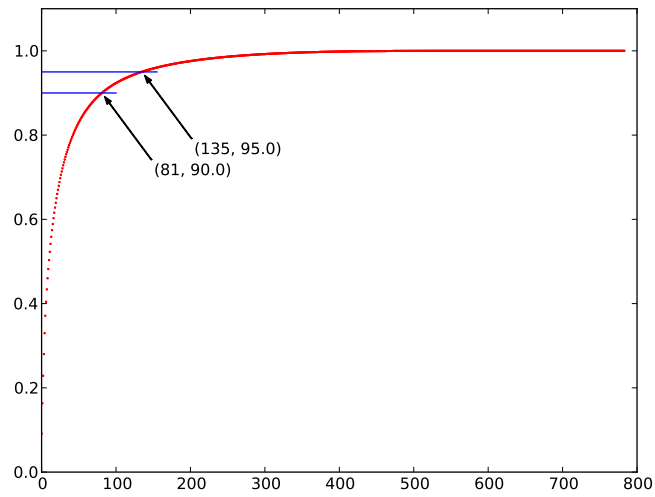


Figure 3: Plot of the normalized eigenvalues

1.3 WU3

Do these look like digits? Should they? Why or why not? (Include the plot in your write-up.)

Although most of the images in figure 4 do not look like digits, a few do resemble simple primitives which some digits share. These eigenvectors likely contribute significantly to the basic structure of some digits, but most of the eigenvectors are encoding detail information which is not easily recognizable. We can show how only a few eigenvectors encode some basic structure by taking the eigenvectors and the projected data and trying to reconstruct the original dataset. With 784 eigenvectors, the reconstructed data appears to be very accurate as shown in 5, but with only 5 we still see that some simple structures have been captured as shown in 6 which shows why some of the eigenvectors in figure 4 may look like simple digit structures.

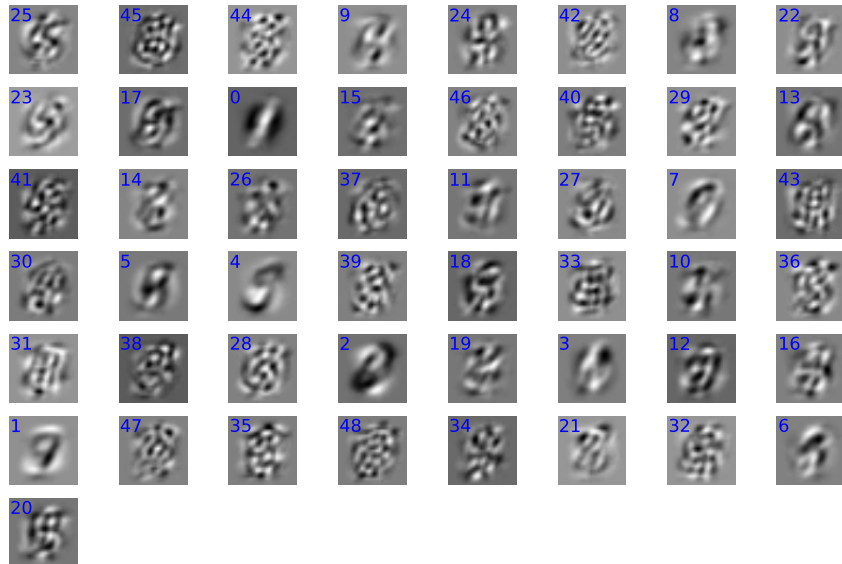


Figure 4: Plot of eigenvectors using vanilla pca

1.4 WU4

Why does vanilla PCA find this data difficult? What is the significance of the relatively large value of the eigenvalues here? Vanilla PCA can only deal with data that is linearly separable. In this dataset, the data is not linearly separable, so vanilla PCA cannot handle it well. The eigenvalues are large



Figure 5: Reconstructed digit images using 784 eigenvectors



Figure 6: Reconstructed digit images using only 5 eigenvectors

in both directions, that is, the variance is large in both directions. Neither of the eigenvectors maximize the variance of the data. This is consistent with the data plot, where the data is in a circle. So none of the directions linear PCA gives us is significantly better than the other.

Vanilla PCA can only deal with the data that is linear separable. In this dataset, the data is not linear separable, so vanilla PCA cannot handle it well. The eigenvalues are large in both directions, that is, the variance is large in both directions. Neither of the eigenvectors maximize the variance of the data. This is consistent with the data plot, where the data is in a circle. So none of the directions linear PCA gives us is significantly better than the other.

Because the original dataset is a sphere shape, so the 2 eigenvalues are almost equal. No matter which eigenvector is chosen, there will be almost half of the information lost. Linear PCA doesn't work well on this data set. Due to the loss of nearly half of the information of the data, after we restore the data, it looks nothing like the original data.

1.5 WU5

Did PCA do what we might want it to? Why or why not? Include the plot to justify your answer.

1.6 WU6

How do the eigenvalues here compare to the linear case? What does this tell you? How does the plot look? How might this be useful for supervised learning?

The eigenvalues separate data into two separate clusters better than the linear case. This is useful for supervised learning because then since we know the labels, we can try different kernels, use the one that gives us the largest margin to make the training data separable, then train discriminative classifiers on the data. It allows us to make linearly inseparable data into separable ones before we feed it to a classifier.

1.7 WU7

Experiment with different kernels, and perhaps interpolations of different kernels. Try to find a kernel that gets as much of the variance on the first two principal components as possible. Report your kernel and a plot of the data projected into 2d under that kernel. *rbf0₂* does well so far.

2 HMMs:Viterbi

2.1 WU8

Find two different observation sequences of length four that differ only in one place, but differ in more than one place in the guessed output.

One example of this effect can be found with observation sequences $[0,1,1,1]$ and $[0,1,2,1]$ which differ only in the third observed emission, but which yield predictions of $[0,0,0,0]$ and $[0,0,1,1]$ respectively. Thus, the predictions differ in two states where the observations only differ in one.

3 HMMs:Forward-backward

3.1 WU9

Why does EM settle in this configuration? What is it saying? What happens when you use three states instead of two? What's the resulting probability of the data? What about four states? What happens and why? EM settles in

this configuration because it is a local minima. The $\hat{\pi} = [1, 0, 0]$ means that we always start at state 0. We also get $P(X_t = 0 | X_{t-1} = 0) = 0$, i.e. when we transition we always go to a different state. All together the results say that states alternate and that is the most likely state sequences.

If we use three states instead of two, similar behavior happens with $\hat{\pi}$. One state has a probability of 1 and the rest gets 0. But in the transition probability, we do get a single state that has positive probabilities for more than one states.

When we increase the states to four, the behavior is still similar in a way that we have very high probabilities for specific transitions and emissions, but the probability of transition is more distributed.

This is because the way we produce the initial HMM is random. When there are only two states, the only way for the final probabilities to be well distributed is if initial transition from state 0 to state 0 and state 0 to state 1 is close i.e. $1/2, 1/2$. If anything, it's very easy for EM to exaggerate the probability of the bigger one.

But as the number of states increases, the likelihood of uniformly distributed in the initial HMM probabilities increases. So we get more reasonable (unskewed) results for EM.

3.2 WU10

Run EM on this data using two states. Run it a few times until you get a final log probability at least -9250. (It should happen in two or three trials. This will take a few minutes to run.) When you do this, it will print out the resulting initial state probabilities, transition probabilities and emission probabilities. What has

it learned? Include all of this output in your write-up. The output of EM on this data is:

Initial state probabilities:

state(0): 7.4252e-202

state(1): 1

Transition probabilities:

FROM\TO

0 1

0 0.245207 0.754793

1 0.720637 0.279363

Emission probabilities:

State 0:

_ 0.362509

e 0.232741

a 0.125778

o 0.120613

i 0.104855

u 0.04369

y 0.00345083

t 0.00258696

k 0.0019353

q 0.00121219

g 0.000467752

s 0.000115461

c 2.34141e-05

p 1.62641e-05

l 3.6812e-06

d 1.16317e-06

b 1.69503e-07

h 1.26829e-08

w 8.21573e-12

r 1.47438e-14

f 5.26943e-15

m 4.67596e-15

j 2.11457e-16

n 2.45545e-17

x 4.24731e-29

v 3.78819e-34

z 3.08129e-51

State 1:

t 0.141537

r	0.102944
n	0.0971604
s	0.0970502
h	0.094847
d	0.0670858
l	0.0636134
c	0.0479795
m	0.0439535
f	0.0347001
y	0.0337207
p	0.0329496
w	0.0306518
b	0.0260249
g	0.0244221
v	0.0202417
k	0.0166601
_	0.0155547
a	0.00316851
x	0.00173501
j	0.00173501
u	0.00168636
z	0.000578336
o	2.52571e-07
e	6.20806e-09
i	2.8181e-11
q	1.18265e-17

It has learned the probability of spaces and vowels vs the probability of consonants! State 0 emits all spaces or vowels and state 1 emits consonants. It also makes sense that the transition from state 0 to state 0 is lower than that of state 1, because it's more likely that consonants are followed by consonants than that vowels are followed by vowels.